

Fourth Question: Working With Real Data

Arnau Abella

03/04/2020

Decathlon Dataset


It is easy to see from the Variables factor map of the PCA that the points is inversly proportional to the rank i.e. the more points you get the lower rank you achieve (lower rank is better)._

From the chart we also see that either the variable has an effect to the rank or to the points, which, at the end of the day, is equivalent. Notice that some results do not have the same impact on the puntuation/rank such as 1500m.

```
# Load the dataset and preprocess it.
```

```
data(decathlon)
colnames(decathlon)[c(1,5,6,10)]<-c("x100m", "x400m", "x110m.hurdle", "x1500m")
colnames(decathlon)
```


```
competition <- which(colnames(decathlon) == "Competition")
plot(decathlon[, -c(competition)])
```




```
# cor(decathlon[, -c(competition)])
```

```
# Remove the dependent variable score.
pca<-PCA(decathlon[, -c(competition)])
```

Individuals factor map (PCA)



Variables factor map (PCA)




Principal Component Regression

From the cumulative percentage of variance, we need at least 4 principal components to have an accumulative variance $\geq \frac{2}{3}$.

```
library(FactoMineR)
competition <- which(colnames(decathlon) == "Competition")

pca$eig
plot(pca$eig[,1], type="o", main="Scree Plot")
```

Scree Plot



```
summary(pca)

decathlon$PC1<-pca$ind$coord[,1]
decathlon$PC2<-pca$ind$coord[,2]
decathlon$PC3<-pca$ind$coord[,3]
decathlon$PC4<-pca$ind$coord[,4]
reg_pc<-lm(Points~PC1 + PC2 + PC3 + PC4, data=decathlon)
```

```
summary(reg_pc)
```

```
##
## Call:
## lm(formula = Points ~ PC1 + PC2 + PC3 + PC4, data = decathlon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -77.522 -35.990   5.294  33.767  89.454
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
##
```

```
## (Intercept) 8005.3659      6.9336 1154.574 < 2e-16 ***
## PC1         152.1889      3.1784  47.882 < 2e-16 ***
## PC2           0.3998      5.2561   0.076 0.93979
## PC3        -17.8926      5.8290  -3.070 0.00406 **
## PC4          41.6555      6.5175   6.391 2.09e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 44.4 on 36 degrees of freedom
## Multiple R-squared:  0.9849, Adjusted R-squared:  0.9832
## F-statistic: 585.7 on 4 and 36 DF,  p-value: < 2.2e-16
```


Testing the assumptions of the regression model

- Normality: the error term does follow a normal distribution which is desired.
- Homogeneity: the variance is homogeneous.
- Independence of errors: the errors do have correlation which may affect the results.

```
# Normality
shapiro.test(residuals(reg_pc))
```

```
##
## Shapiro-Wilk normality test
##
## data:  residuals(reg_pc)
## W = 0.96366, p-value = 0.2108
```

```
# Homogeneity
plot(residuals(reg_pc))
```



```
bptest(reg_pc)
```

```
##
## studentized Breusch-Pagan test
##
```

```
## data: reg_pc
## BP = 13.223, df = 4, p-value = 0.01024
# Independence of errors
dwtest(reg_pc, alternative = "two.sided")

##
## Durbin-Watson test
##
## data: reg_pc
## DW = 1.0195, p-value = 0.0004396
## alternative hypothesis: true autocorrelation is not 0
```

Predicting the points of an athlete using the regression model

Nota bene, the RMSE is small compared to the points scale so we can conclude that the model is accurate “enough”.

```
n <- nrow(decathlon)
train.sample <- sample(1:n, round(0.67*n))

train.set <- decathlon[train.sample, ]
test.set <- decathlon[-train.sample, ]

train.model <- lm(Points ~ PC1+PC2+PC3+PC4 , data = train.set)
summary(train.model)

##
## Call:
## lm(formula = Points ~ PC1 + PC2 + PC3 + PC4, data = train.set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -60.83  -33.70  -13.57   34.43   65.69
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7998.1407     8.5430  936.223  < 2e-16 ***
## PC1           152.0412     4.4297   34.323  < 2e-16 ***
## PC2             0.5286     8.0539    0.066  0.94826
## PC3          -25.6946     7.0097   -3.666  0.00136 **
## PC4            42.7321     7.4789    5.714  9.54e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 43.97 on 22 degrees of freedom
## Multiple R-squared:  0.9826, Adjusted R-squared:  0.9795
## F-statistic: 311 on 4 and 22 DF, p-value: < 2.2e-16

yhat<-predict(train.model, test.set, interval="prediction")
yhat

##              fit      lwr      upr
## MARTINEAU 7776.188 7680.956 7871.420
## NOOL       7708.842 7604.981 7812.704
## Sebrle     8851.531 8741.777 8961.286
```

```
## Karpov      8702.490 8594.922 8810.058
## Warners     8338.320 8236.283 8440.358
## Zsivoczky   8298.769 8199.168 8398.369
## Schwarzl    8047.820 7950.861 8144.778
## Smith       8055.547 7954.041 8157.053
## Ojaniemi    8035.316 7941.154 8129.478
## Qi          7902.937 7807.772 7998.101
## Drews       7838.609 7729.806 7947.412
## Terek       7779.868 7682.123 7877.613
## Lorenzo     7582.013 7478.281 7685.745
## Casarsa     7366.519 7248.579 7484.459
```

```
y<-test.set$score

error<-cbind(yhat[,1,drop=FALSE],y,(yhat[,1]-y)^2)
sqr_err<-error[,1]
sse<-sum(sqr_err)

# Root Mean Square Error = sqrt(SSE/N)
RMSE<-sqrt(sse/(nrow(test.set)))
RMSE
```

```
## [1] 89.55635
```